

# CS 383: Machine Learning

Prof Adam Poliak

Fall 2024

09/11/2024

Lecture 03

# Outline

- **Featurization & K-NNs**
- Decision Trees
- Entropy

# Predicting Graduation on Time

Features:

- Major:
  - Computer Science, Art History, English
- Dorm:
  - Rhoads, Pembroke, Merion

Major:	{0,1,2}
Dorm:	{0,1,2}

Given a new student, how can we compute distance between her and the training examples if features are categorical?

# Distance metric - are all features equal?


Day	Outlook	Temperature	Humidity	Wind	PlayTennis ( $y$ )
$x_1$	Sunny	Hot	High	Weak	No
$x_2$	Sunny	Hot	High	Strong	No
$x_3$	Overcast	Hot	High	Weak	Yes
$x_4$	Rain	Mild	High	Weak	Yes
$x_5$	Rain	Cool	Normal	Weak	Yes
$x_6$	Rain	Cool	Normal	Strong	No
$x_7$	Overcast	Cool	Normal	Strong	Yes
$x_8$	Sunny	Mild	High	Weak	No
$x_9$	Sunny	Cool	Normal	Weak	Yes
$x_{10}$	Rain	Mild	Normal	Weak	Yes

*Data from Machine Learning by Tom Mitchell (Table 3.2)*

# Outline

- Featurization & K-NNs
- **Decision Trees - Overview**
- ID3 Decision Tree Algorithm

# Decision Tree example (tennis)




Each internal node: test one feature

Each branch from node: selects one value of the feature

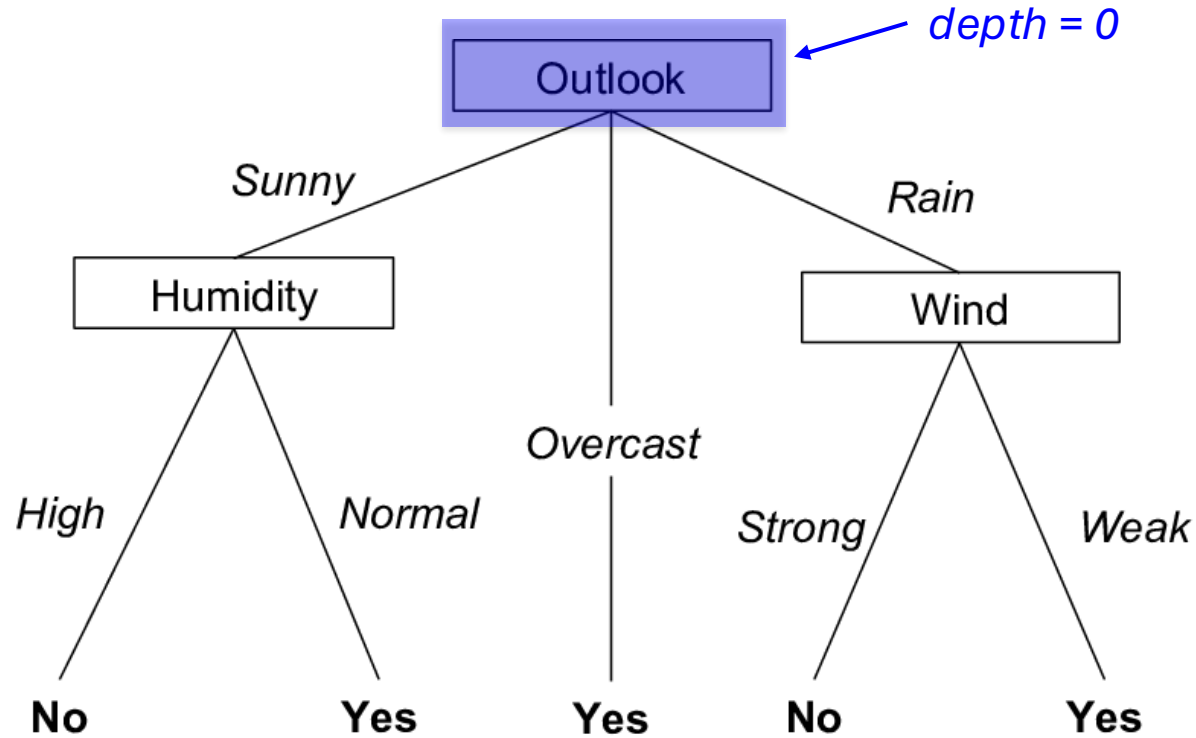
Each leaf node: predict  $y$

# Decision Tree example (tennis)



Key term: *depth*

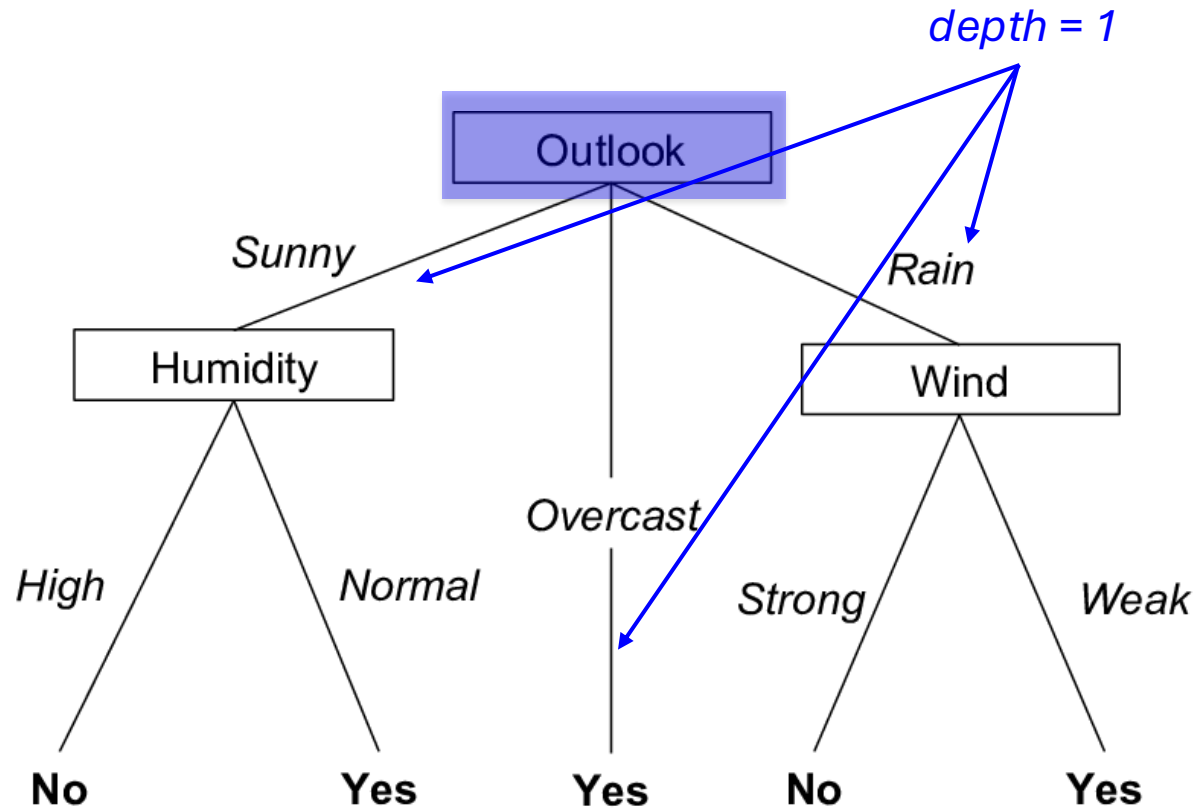
# Decision Tree example (tennis)



Key term: *depth*




# Decision Tree example (tennis)



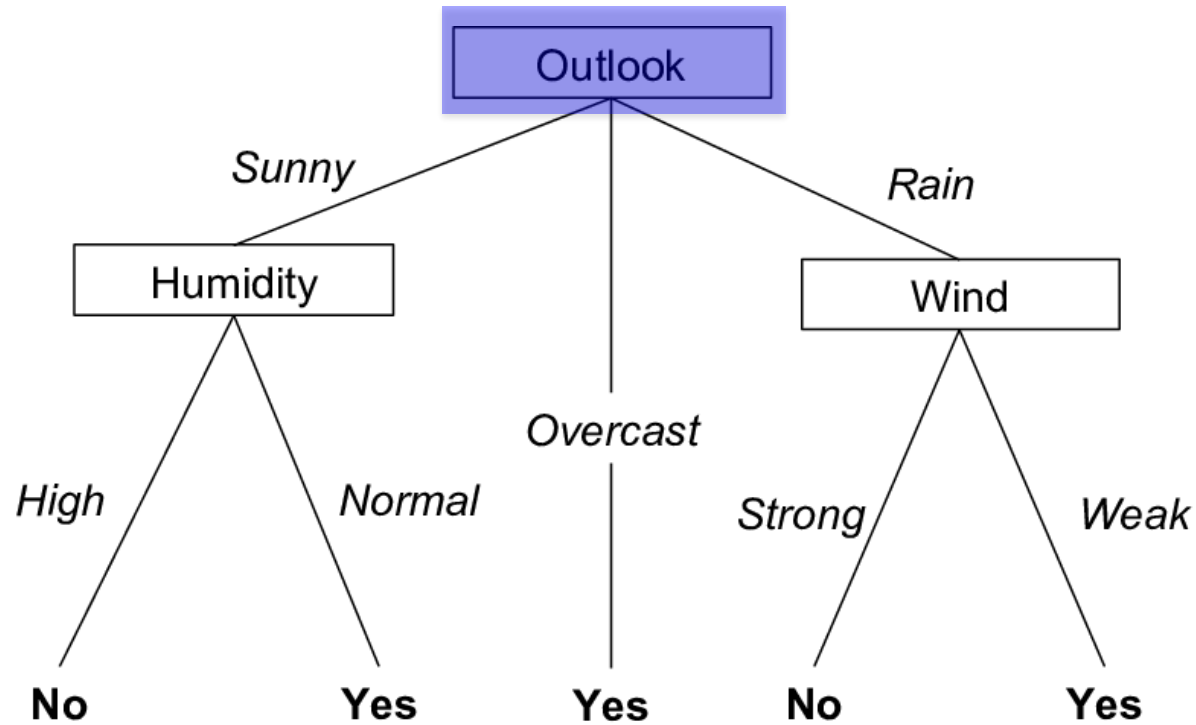
Key term: *depth*

# Decision Tree example (tennis)



Key term: *depth*


# Decision Tree example (tennis)



(test example)  $x =$

Outlook	Temp	Humidity	Wind
Rain	Hot	High	Strong

# Decision Tree example (tennis)




(test example)  $x =$


Outlook	Temp	Humidity	Wind
Rain	Hot	High	Strong

$\hat{y} = \text{No}$

# Continuous Features




# Continuous Features




$$x_2 \leq 3$$

# Continuous Features



$$x_2 \leq 3$$

# Continuous Features





# Decision Trees: Pros vs Cons

Discuss with a partner! Think about:

- \* training and testing
- \* featurization
- \* runtime
- \* human factors

# Decision Trees: Pros vs Cons

- Very interpretable! Easy to say *why* we made a classification (can point to which features)
- Compact representation and fast predictions
- Can be brittle (not looking at each example holistically)
- Featurization and implementation difficulties

# Outline

- Featurization & K-NNs
- Decision Trees - Overview
- **ID3 Decision Tree Algorithm**

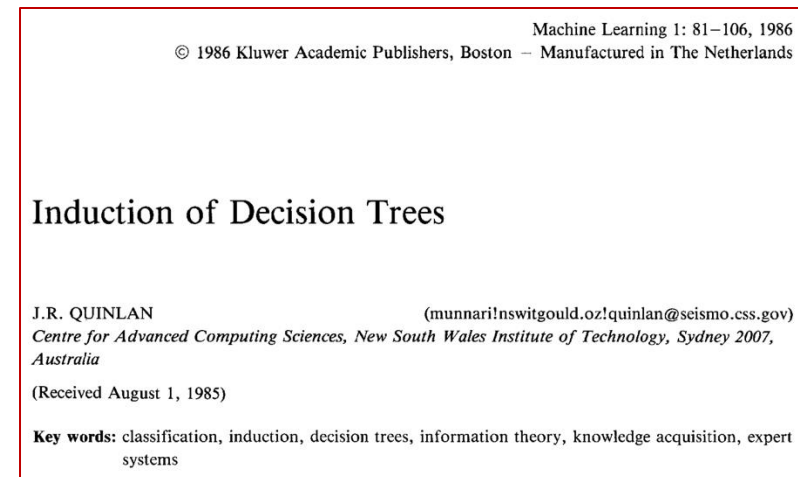
# ID3 Decision Tree Algorithm

Select feature that “best” informs label prediction (i.e.  $y$ )

**Divide:** partition data into branches based on their value at this feature

**Conquer:** recurse on each partition

Posted as optional reading



# Top-Down Decision Tree Algorithm


**MakeSubtree**( $D, F$ )

- if stopping criteria met


  - make a leaf node  $N$

  - determine class label/probabilities for  $N$

# Top-Down Decision Tree Algorithm

 Dataset (X,y)  
**MakeSubtree**( $D$ ,  $F$ )  
if stopping criteria met  
    make a leaf node  $N$   
    determine class label/probabilities for  $N$   
    .

# Top-Down Decision Tree Algorithm

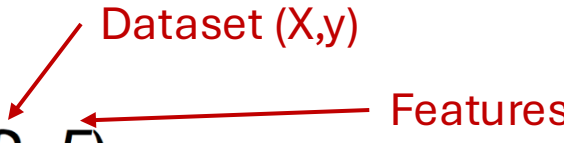
  
`MakeSubtree( $D$ ,  $F$ )`

if stopping criteria met

make a leaf node  $N$

determine class label/probabilities for  $N$

# Top-Down Decision Tree Algorithm



```
MakeSubtree( $D$ ,  $F$ )  
  if stopping criteria met  
    make a leaf node  $N$   
    determine class label/probabilities for  $N$   
  else  
    make an internal node  $N$   
     $S = \text{FindBestFeature}(D, F)$   
    for each outcome  $k$  of  $S$   
       $D_k =$  subset of instances that have outcome  $k$   
       $N.\text{child}[k] = \text{MakeSubtree}(D_k, F - S)$   
  return subtree rooted at  $N$ 
```



# Top-Down Decision Tree Algorithm

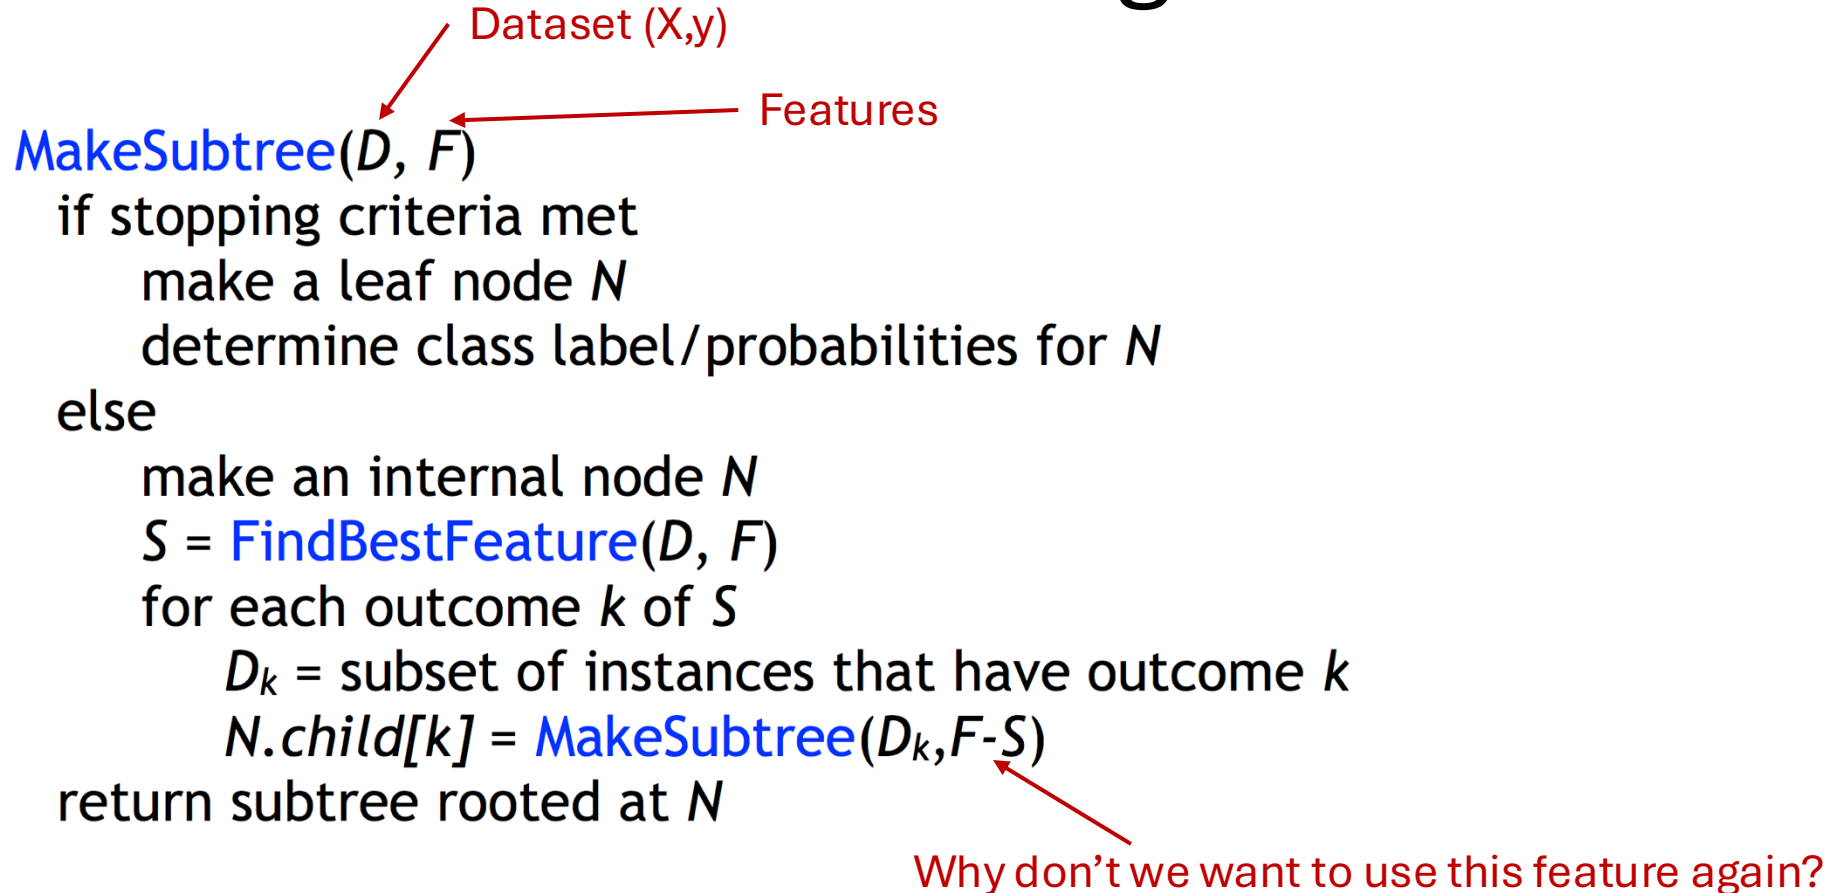
```
MakeSubtree( $D, F$ )  
  if stopping criteria met  
    make a leaf node  $N$   
    determine class label/probabilities for  $N$   
  else  
    make an internal node  $N$   
     $S = \text{FindBestFeature}(D, F)$   
    for each outcome  $k$  of  $S$   
       $D_k =$  subset of instances that have outcome  $k$   
       $N.\text{child}[k] = \text{MakeSubtree}(D_k, F-S)$   
  return subtree rooted at  $N$ 
```

Dataset ( $X, y$ )

Features

Why don't we want to use this feature again?

# Top-Down Decision Tree Algorithm

  
**MakeSubtree**( $D, F$ )  
if stopping criteria met  
    make a leaf node  $N$   
    determine class label/probabilities for  $N$   
else  
    make an internal node  $N$   
     $S = \text{FindBestFeature}(D, F)$   
    for each outcome  $k$  of  $S$   
         $D_k =$  subset of instances that have outcome  $k$   
         $N.\text{child}[k] = \text{MakeSubtree}(D_k, F-S)$   
return subtree rooted at  $N$   
  
Why don't we want to use this feature again?

**Now: Handout 3 + think about:  
what design choices do we need to make?**

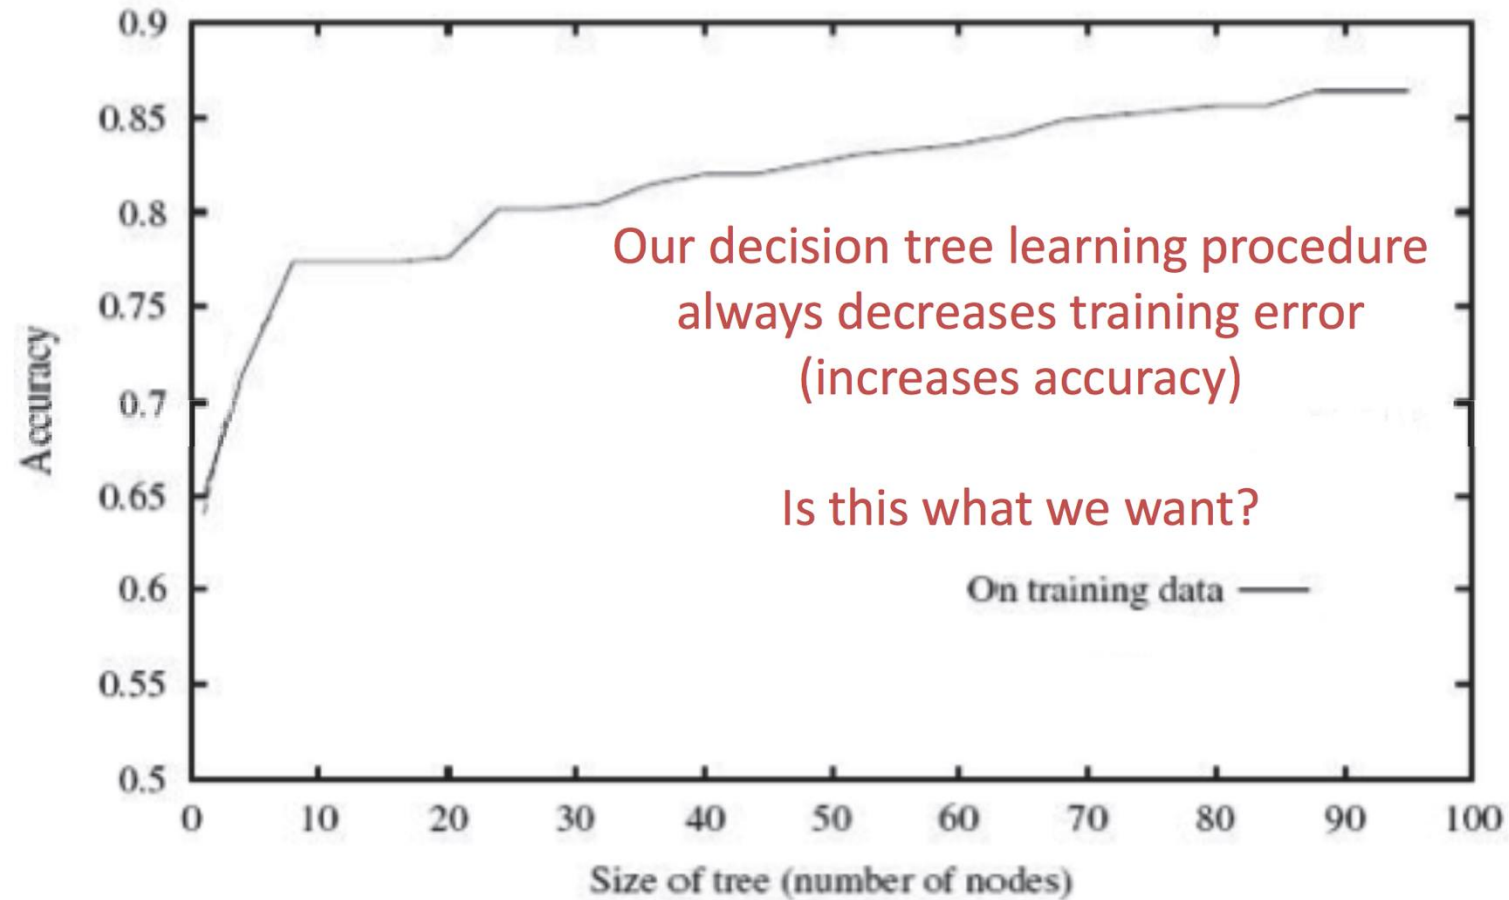
# Design choice: stopping criteria

1. All the data points in our partition have the same label
2. No more features remain to split on
3. No features are informative about the label
4. Reached (user specified) max depth in the tree


# Next class

- “Best feature”
- HW01 due Friday night
- Thursday lecture back on schedule
  - Reading quiz on Thursday

# Overfitting



# Overfitting



# Overfitting definition

Consider a hypothesis (tree):  $h$

- Training error:  $error_{train}(h)$
- Error over all possible data:  $error_D(h)$

# Overfitting definition

Consider a hypothesis (tree):  $h$

- Training error:  $error_{train}(h)$
- Error over all possible data:  $error_D(h)$

A hypothesis  $h$  **overfits** training data if there exists another hypothesis  $h'$  s.t.

$$error_{train}(h) < error_{train}(h') \text{ AND } error_D(h) > error_D(h')$$



# Avoiding overfitting in decision trees

- Stop when leaf label reaches a certain fraction (i.e. 95% “yes”, 5% “no”)
- Set a maximum depth for the tree
- Set a minimum number of examples in leaf (i.e. if we have a 2-1 split, stop)

HW2 implementation

