

# CS 383: Machine Learning

Prof Adam Poliak

Fall 2024

09/24/2024

Lecture 08

# Announcements

HW02 is due tonight night

HW03 is due next Tuesday night

- **Reading quiz: Thursday**
  - Duame 9.3 (2 pages)

# Proposed updated schedule

Midterm 1 was Thursday October 3<sup>rd</sup>

3 lectures this week

lecture on Wednesday 10/02 but no lecture on Thursday 10/03 (was supposed to be midterm 1)

No lecture Wednesday 10/09

HW02 decision trees due tonight, HW03 polynomial regression due next Tuesday 10/01, HW04 naive Bayes due Tuesday 10/08 (it'll be a shorter assignment)


Midterm 1 on Thursday 10/10

# Outline



Normal equations solution

Regularization

# Cost Function



# Cost Function



$$J(0.5) = \frac{1}{2} [(0.5 - 1)^2 + (1 - 2)^2 + (1.5 - 3)^2] = 1.75$$

# Outline

**Normal equations solution**

Regularization

# Assumptions of linear regression

Explanatory (independent) and response (dependent) variables have a linear relationship

Instances are independent of each other

Residuals have a normal distribution with mean 0

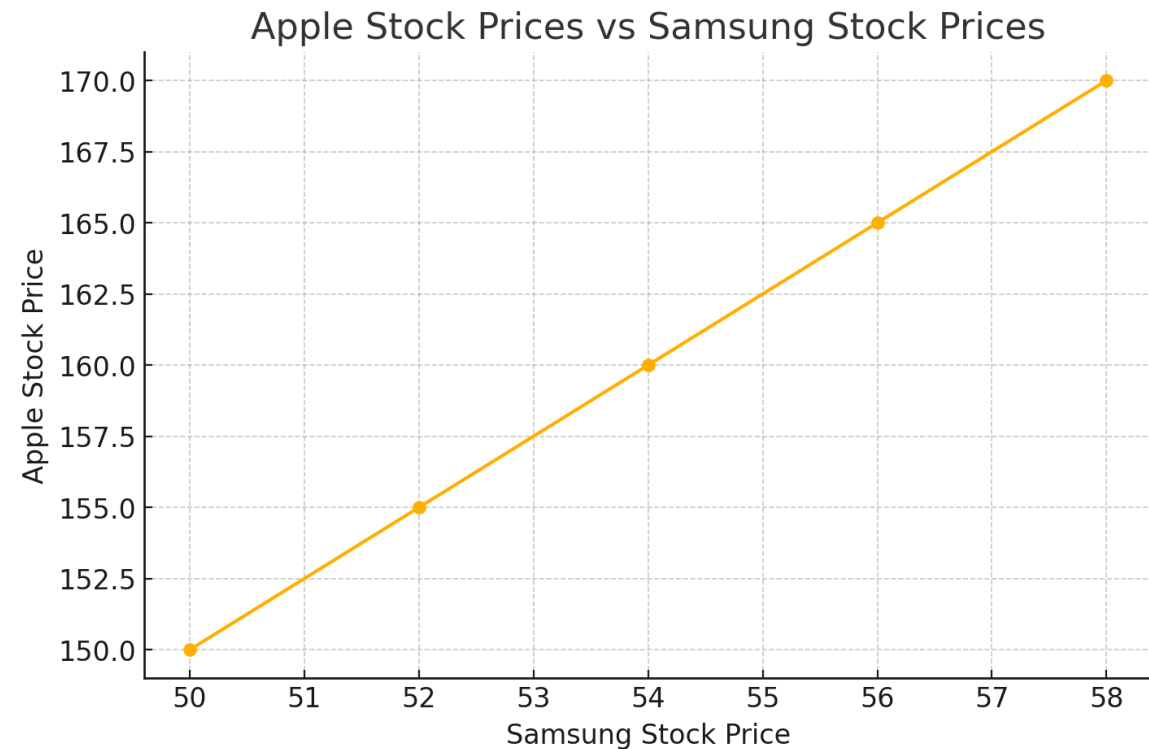
The variance of the residuals is the same for an  $X$  (independent variable) - Homoscedacity



# Linear Regression example with ChatGPT

<https://chatgpt.com/share/66f2e089-95a0-8011-a567-7e75bd93261c>


5 examples predicting  
Apple stock from  
Samsung stock



# Linear Regression example with ChatGPT

<https://chatgpt.com/share/66f2e089-95a0-8011-a567-7e75bd93261c>

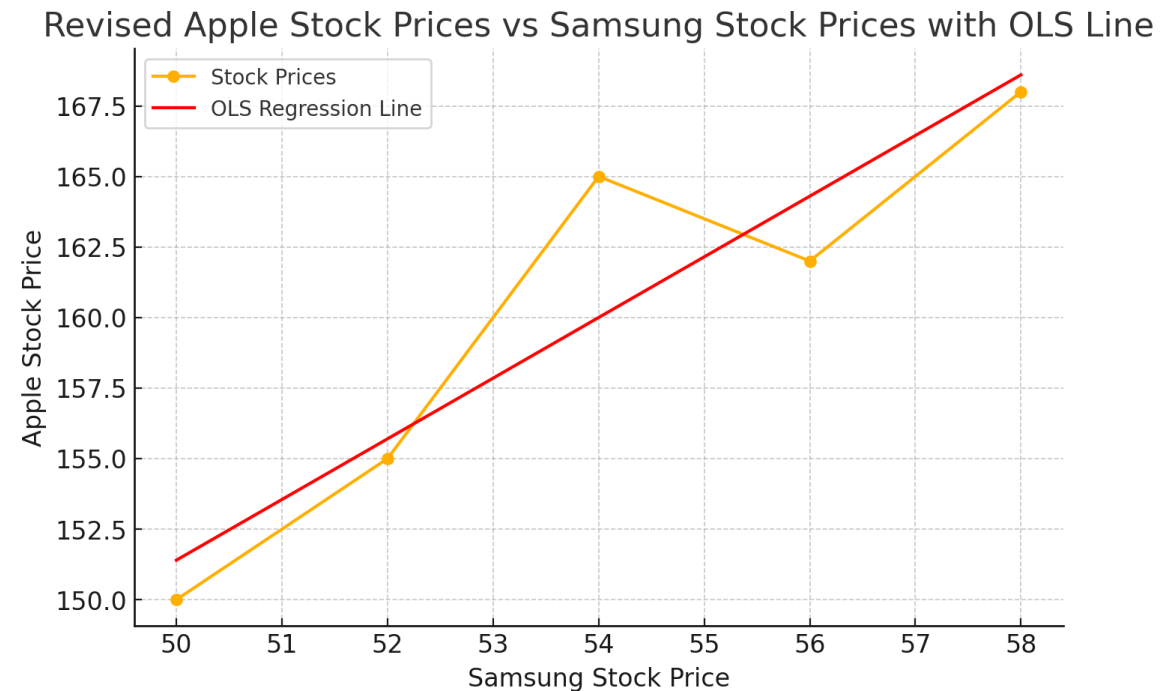
5 revised examples  
predicting  
Apple stock from  
Samsung stock



# Linear Regression example with ChatGPT

<https://chatgpt.com/share/66f2e089-95a0-8011-a567-7e75bd93261c>


5 revised examples  
predicting  
Apple stock from  
Samsung stock



# Linear Regression example with ChatGPT

<https://chatgpt.com/share/66f2e089-95a0-8011-a567-7e75bd93261c>


5 revised examples  
predicting  
Apple stock from  
Samsung stock



# Linear Regression example with ChatGPT

<https://chatgpt.com/share/66f2e089-95a0-8011-a567-7e75bd93261c>


1000 examples  
predicting  
Apple stock from  
Samsung stock



# Linear Regression example with ChatGPT

<https://chatgpt.com/share/66f2e089-95a0-8011-a567-7e75bd93261c>

Residuals and feature  
are uncorrelated




Proof: <https://statproofbook.github.io/P/slr-rescorr.html>

# Linear Regression example with ChatGPT

<https://chatgpt.com/share/66f2e089-95a0-8011-a567-7e75bd93261c>


Distribution of residuals is normal, with mean 0



# Linear Regression example with ChatGPT

<https://chatgpt.com/share/66f2e089-95a0-8011-a567-7e75bd93261c>

Predicting (red line)  
Apple price from  
Samsung price where  
variables are  
uncorrelated






# Linear Regression example with ChatGPT

<https://chatgpt.com/share/66f2e089-95a0-8011-a567-7e75bd93261c>

Distribution of residuals is not normal when response and predictor variable are uncorrelated



# Normal Equation

$$J(\vec{w}) = \frac{1}{2} \sum_i^n \sum_k^p (w_0 + w_1 x_{i,1} + w_2 x_{i,2} + \dots + w_p x_{i,p} - y_i)^2$$

$$J(\vec{w}) = \frac{1}{2} (\vec{w}X - y)^2$$

# Normal Equation

$$J(\vec{w}) = \frac{1}{2} \sum_i^n (w_0 + w_1 x_i - y_i)^2$$

Now we have more than 1 feature:

$$J(\vec{w}) = \frac{1}{2} \sum_i^n \sum_k^p (w_0 + w_1 x_{i,1} + w_2 x_{i,2} + \dots + w_p x_{i,p} - y_i)^2$$

How many partial derivatives do we need to compute?

# Linear Algebra Review

- Transpose
- $A^T$  swap all columns and rows

$$\begin{aligned}(Ax)^T &= \\ &= x^T A^T\end{aligned}$$

# Pros and Cons

## Gradient Descent

- Requires multiple iterations
- Need to choose  $\eta$
- Works well when  $n$  is large
- Can support online learning



## Normal Equation

- Non-iterative
- No need to choose  $\eta$
- Slow if  $p$  is large
  - Matrix inversion is  $O(p^3)$

# Toy example, iteration 1


This is what you  
should have  
obtained in Handout  
7!

iteration: 1, cost: 0.410000




# Toy example, iteration 2

iteration: 2, cost: 0.350001



# Toy example, iteration 12


iteration: 12, cost: 0.138047







# Toy example, iteration 40

iteration: 40, cost: 0.014064



# Toy example, iteration 100

iteration: 100, cost: 0.000105






# Outline

Normal equations solution

**Regularization**


# Generalization Error

Example: price vs. size (i.e. of a house or car)




# Generalization Error


Example: price vs. size (i.e. of a house or car)



underfitting  
(high bias)



correct fit



overfitting  
(high variance)

# Generalization Error

## Structural error:

Hypothesis space cannot model true relationship

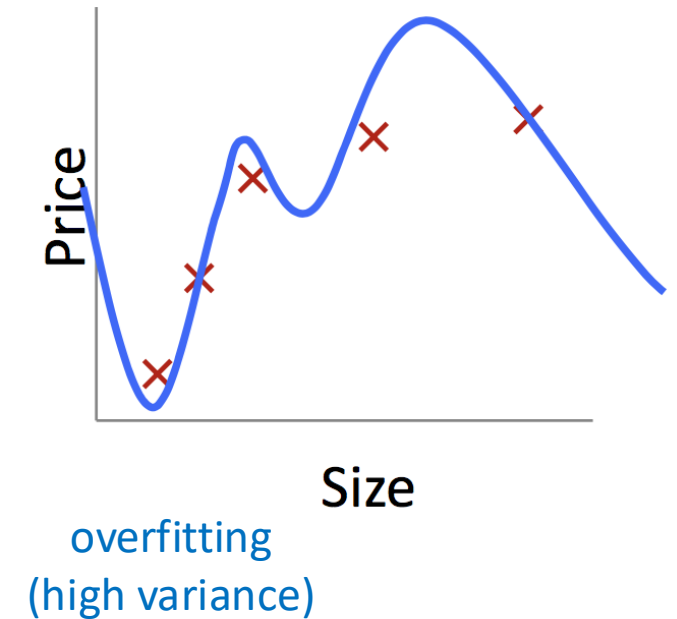
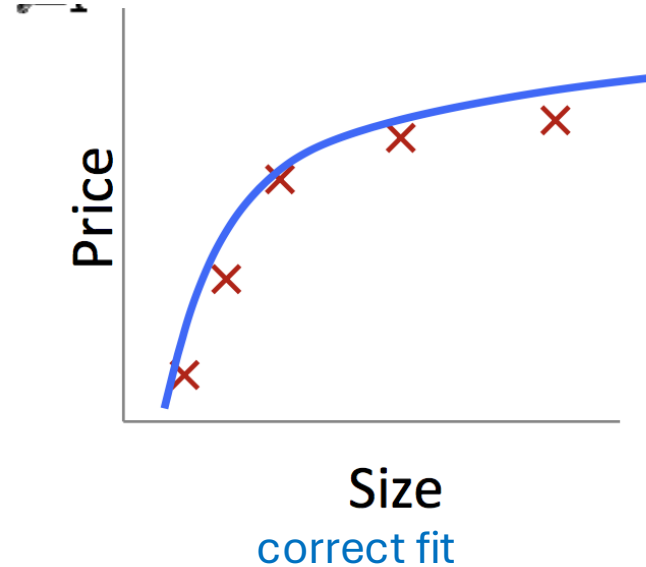

- More data doesn't help
- Need a more flexible model

## Estimation (approximation) error:

Hypothesis space *can* model true relationship, BUT hard to identify correct model due to large hypothesis space, small  $n$ , or noise

- ☐ Reduce hypothesis space
- ☐ Add more data

balance  
↔



# Regularization

What if ...

- we have a limited # of training examples ( $n < p$ ), or
- we want to automatically control the complexity of the learned hypothesis?

# Regularization

What if ...

- we have a limited # of training examples ( $n < p$ ), or
- we want to automatically control the complexity of the learned hypothesis?

Idea: penalize large values of  $w_j$

Why prefer small weights?



# Regularization

What if ...

- we have a limited # of training examples ( $n < p$ ), or
- we want to automatically control the complexity of the learned hypothesis?

Idea: penalize large values of  $w_j$

Why prefer small weights?

- if large weights, small change in feature can result in large change in prediction
- prevent giving too much weight to any one feature
- might prefer zero weight for useless features

# Common Regularizers

$$||\vec{w}||_0 = \sum_{j:w_j \neq 0} 1$$

$L_0$  norm

- Number of non-zero entries
- Minimizing  $L_0$  norm is NP hard

$$||\vec{w}||_1 = \sum_{j=1}^p |w_j|$$

$L_1$  norm

- Sum of magnitude of weights
- Not differentiable

$$||\vec{w}||_2 = \sqrt{\sum_{j=1}^p w_j^2}$$

$L_2$  norm

- Sum of squared weights
- Differentiable